

# Dmitry Prom

AI Engineer & Data Analyst (LLM-based analytics)

OPEN TO SENIOR / LEAD APPLIED AI · REMOTE · UTC+3 · OPEN TO TEAMS ABROAD

xpromx94@yandex.ru  
github.com/94spec  
94spec.github.io  
leadchat.online

## SUMMARY

I build production LLM systems that take manual work off revenue teams: a customer conversation becomes verifiable data. And I do the analysis on top of that data – the analysis behind decisions on sales, marketing, product and pricing.

I own the whole path: business policy → model behaviour → architecture → implementation → evaluation → CRM and BI integration → cost and latency → the report people act on. I do not manage people: I work alongside developers and own the AI loop myself.

## KEY RESULTS · EVERY NUMBER WITH ITS DENOMINATOR

|                                                                                                                                                                                      |                                                                                                                                             |                                                                                                                                                            |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>100% / ≤20%</b><br>of eligible conversations evaluated, instead of manual sampling<br>denominator – conversations passing the duration and content filter                         | <b>98%</b><br>agreement between AI and expert decisions<br>denominator – 500 deals in a frozen golden set; critical recall gated separately | <b>–88%</b><br>cost of the key LLM scenario per 1,000 deals<br>with output schema, benchmark and quality gates unchanged                                   |
| <b>45–50% · ×1,2</b><br>of routine consultations without an operator; qualified-lead lift<br>denominator – eligible routine consultations; leads – month one vs the previous process | <b>≥ 1 млн ₽</b><br>monthly contact-centre saving<br>cost of servicing the same consultation flow, before vs after automation               | <b>90 / 464 / 246</b><br>production pipelines, checklist criteria, summary fields and detectors<br>plus 55 active checklists and 32 summary configurations |

AI quality control across every customer communication

MAJOR EDTECH PLATFORM (RU) · BUILT THE LOOP END TO END

Quality control used to listen to a manual sample of calls. Now the communication flow is transcribed, filtered, scored against checklists and turned into structured data for managers, CRM and BI – covering every employee who talks to customers. Mine: platform logic, system prompts, model selection, pipelines, scoring rules, delivery and calibration on live traffic.

Evaluation stopped being a sample: 100% of eligible conversations instead of ≤20%.

Analytics that drives decisions

SALES · MARKETING · PRODUCT AND PRICING · COMPETITORS

Tasks arrive from the business as tickets with a goal and an economic rationale. I build the cut, validate the sample and deliver an answer that can be acted on: what is happening, what supports it and what the data cannot show. Studies run up to 3,000 deals and up to a six-month horizon, across calls and chats.

- **Why conversion dropped over a quarter.** The month-over-month view showed the problem was not the sales team: traffic was cooling, purchase readiness was falling, money objections were growing. → decisions on traffic buying, script and pricing pitch.
- **"Too expensive" is not always about price.** A two-level reason taxonomy: 'cost' at the top, but inside the group the dominant driver is unavailable instalment financing. That is a payment-instrument problem, not a pricing one.
- **Part of the sales team's load is not sales.** Classifying inbound requests by meaning and by owner surfaced service load inside the sales funnel and a conversion denominator that measured the wrong thing.

Voice AI

VOICE AGENT INSTEAD OF AN OUTSOURCED CONTACT CENTRE

I turned the contact-centre policy into a live stateful dialogue: barge-in, clarifications, script branching, answers from the knowledge base. The agent replies instantly, works around the clock and never forgets product details.

45-50% of routine consultations without an operator, ×1.2 qualified leads in month one, saving from 1M RUB per month.

Enterprise RAG

TWO INTERNAL ASSISTANTS · 100+ DAILY USERS

Assistants for sales and corporate documentation. I owned chunking, retrieval, context assembly, grounding and answer quality. The real work is not the model but document preparation – and making the assistant honestly say 'this is not in the documents'.

An answer with a source link in under 10 seconds instead of roughly five minutes of manual search.

LLM evaluation and model economics

4 TASK-SPECIFIC GOLDEN SETS · REGRESSION GATES

Separate sets and metrics for scoring, summaries, semantic analytics and violations. Human calibration, regressions and error analysis after every model change. I moved the same validated output from an expensive model to a compact one: long context and format stability break first, so the release is gated on schema validity and recall on critical violations rather than average accuracy.

98% decision agreement on a frozen 500-deal set; -88% cost of the key scenario.

Lead Chat – my own SaaS

LEADCHAT.ONLINE · PAY PER RESULT: 50 ₪ PER CREATED LEAD

A working product, not a demo: the agent runs a goal-driven dialogue, qualifies the customer, handles objections, answers from the knowledge base and takes actions – books, captures the contact, pushes the lead into CRM. Mine: product concept, AI logic and core architecture. A development partner works with me on the technical side. Next.js, Node.js, tRPC, Prisma, PostgreSQL / pgvector, RAG, tool calling, background jobs, billing.

ENGINEERING DECISIONS AND THEIR PRICE

|                   |                                                                                                                                                                                                                                          |
|-------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Transcription     | Managed STT instead of self-hosted Whisper: quality on spontaneous phone speech and diarisation out of the box, predictable price, zero operations. Revisit when volume makes the price gap exceed the cost of operations.               |
| Retrieval         | pgvector in the same PostgreSQL instead of a dedicated vector database: search next to the business filters, one query instead of two stores and their sync. Threshold – when retrieval time becomes noticeable next to model time.      |
| Evaluation        | An LLM judge against frozen expert labels instead of training a classifier per criterion: criteria change with the policy, and the judge explains its decision. The price is drift when the model version changes, so the set is frozen. |
| Human in the loop | Where the cost of a mistake is high, the decision stays with a human. A deliberate limit, not an unfinished feature.                                                                                                                     |

HOW I WORK WITH DATA

- Sample control: how many records were included, how many were lost and why – before any percentages.
  - A denominator for every share, and the same one when periods are compared.
- Observation and explanation are kept apart: a hypothesis is called a hypothesis and comes with a way to test it.
  - Every category is backed by a customer quote, and the stakeholder can verify the labelling by hand.
  - A limitations section: what the data does not show and which conclusions it cannot support.

STACK

IN PRODUCTION, UNDER MY RESPONSIBILITY

Python · TypeScript · FastAPI · Node.js · PostgreSQL · pgvector · SQL · Docker · OpenAI API · Anthropic API · GigaChat · YandexGPT · Dify · RAG · AI Agents · Structured Outputs · Tool Calling · LLM-as-a-Judge · Voice AI · CRM / BI · Git

WORKED WITH, OUTSIDE THAT LOOP

LangChain · LangGraph · Qdrant · vLLM · Hugging Face · Fine-tuning / LoRA · RAGAS · A/B Testing · Guardrails · Prompt Injection · NLP

PUBLIC CODE

|                               |                                                                                                                                                                                                                                            |
|-------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| llm-evaluation-harness        | Dependency-free LLM evaluation framework: four task families, typed error taxonomy, bootstrap intervals, release gates. <a href="https://github.com/94spec/llm-evaluation-harness">github.com/94spec/llm-evaluation-harness</a>            |
| conversation-intelligence-lab | The same analytics methodology on synthetic data: strict output contract, traceable aggregates, reproducible report. <a href="https://github.com/94spec/conversation-intelligence-lab">github.com/94spec/conversation-intelligence-lab</a> |

Employer system names and internal figures are not disclosed. Every figure here comes with its sample, denominator and comparison rule – happy to walk through the details on a call. Current version of this CV: [94spec.github.io/cv/](https://94spec.github.io/cv/)